

Cogenetic Minkowski Functional

Version 1.0.1

Morgan Petrik

Militant.AI

23 June 2026

DOI: [10.5281/zenodo.20808466](https://doi.org/10.5281/zenodo.20808466)

Abstract

Following *Cogenesis* [?]: relation is prior to being; spacetime, metrics, and measurement are **projections** of a deeper process-field, not primitives. Minkowski’s fusion of time with space into a (3+1) manifold was a **modern calculus utilitarian simplification**—an accounting convention for Lorentz covariance and perturbation theory. Standard practice **reified** that convention into “spacetime.” That is the error: fundamentalising time as a fourth dimension on par with space, and treating $(\mathcal{M}, \eta)$ as the relational dynamics of reality rather than as projected bookkeeping.

Cogenetic Minkowski Functional (CMF) is the Minkowski-stage observation functional in that residual-aware framework. It keeps the calculus machinery—flat η , light cones, variational methods—while refusing the reification. η is the **zeroth-order low-residual projection**, not ground truth. Environment enters as C —the environmental density gradient along the observer’s throughline. **Process time** t indexes resolution $R(t)$ —not a rename of t as resolution, but the map $t \mapsto R(t)$ at which the functional is evaluated and sharpens as the process runs.

CMF is a **refinement algorithm** on the **projection calculus**—against the **process calculus** beneath it: project, measure residual, recover flat η only where $|r_{C,R}|$ is negligible. Flat Minkowski remains *usable* when the utilitarian simplification suffices—when ignoring C and $R(t)$ is adequate for the accounting task, not because spacetime is true.

This document specifies the construction, contrasts it with the bolted-on input pattern of standard practice, and separates what Minkowski *stipulates* from what CMF *determines* at each process time and *predicts* over the running observation. A reference implementation with a satellite point-detector example is included.

1 Problem statement

1.1 What the Minkowski calculus is for

On $\mathcal{M} = \mathbb{R}^{3,1}$ with metric $\eta_{\mu\nu} = \text{diag}(-c^2, 1, 1, 1)$, the interval

$$ds^2 = \eta_{\mu\nu} dx^\mu dx^\nu \tag{1}$$

is a **utilitarian simplification**: a flat stage for special relativity, QFT regularisation, and linearised gravity. The apparatus is proven and cheap to use. None of that makes $(\mathcal{M}, \eta)$ a reification of relational reality.

1.2 Where reification goes wrong

Minkowski’s (3+1) fusion was an accounting hack—time promoted to a fourth coordinate so Lorentz transformations act linearly on a single object. Standard practice then:

1. fixes a functional early—e.g. $F[\Phi] = \int |\Phi|^2 d\mu$, an action $S[\Phi]$, or a Green’s function;
2. treats *environment* and *process* as *external inputs* bolted on afterward;
3. treats *events* as primitive points on $\mathcal{M}$;
4. identifies this *calculus projection* with *physical reality*.

Spacetime as primitive is a category error. It hides what any finite observation already carries: observer location, environment along the throughline, running process time t and resolution $R(t)$, and the residual lost when a continuous field is projected to an observable. CMF makes these explicit *inside* the functional, not as afterthoughts.

1.3 Design principle

Process calculus carries the relational dynamics: observer throughline, environment C , process time t , resolution $R(t)$, and the residual $r_{C,R}$ lost on projection.

Projection calculus is the utilitarian simplification Minkowski built from that dynamics: flat $(\mathcal{M}, \eta)$, intervals, Lorentz covariance, variational machinery [?]. It is bookkeeping that converges when $|r_{C,R}|$ is negligible—not a second ontology.

CMF uses both without collapsing them:

Process calculus is primary; $(\mathcal{M}, \eta)$ is projection calculus, not relational reality.

(2)

2 Formal construction

2.1 Base stage

Retain the Minkowski background

$$\mathcal{M} = \mathbb{R}^{3,1}, \quad \eta_{\mu\nu} = \text{diag}(-c^2, 1, 1, 1) \quad (3)$$

as the working stage of **projection calculus** (§??)—not **process calculus**.

2.2 Process field

Introduce a smooth process field

$$\Phi : \mathcal{M} \rightarrow V \quad (4)$$

where V is a finite-dimensional vector space (scalar, spinor, or internal coherence space).

2.3 Context C

C is the declared local condition under which process becomes distinguishable:

$$C = (U, e_a^\mu, A, B, I, E, G) \quad (5)$$

Operational form (primary). For a point detector on a throughline—e.g. a satellite sighting a surface target— C is the environmental density gradient along the observer ray:

$$C(x) = \hat{n} \cdot \nabla \rho_{\text{env}}(x) \quad (6)$$

where ρ_{env} is environmental density (atmosphere, medium, noise background) and $\hat{n}$ is the unit vector from observer to target. C is not metadata bolted onto a fixed functional; it shapes the map itself.

Table 1: Components of context C .

Component	Meaning
$U \subset \mathcal{M}$	Region of distinguishability
e_a^μ	Observer / tetrad frame
A	Apparatus or observation functional
B	Boundary / initial data
I	Preparation / input condition
E	Environmental coupling
G	Gauge / coordinate convention

2.4 Process time and resolution $R(t)$

In the core formalism, R denotes **resolution**—finite distinction capacity—not calendar time. CMF does *not* identify R with coordinate time t . Instead, **process time** t is the running index of the observation process, and t maps to an evolving resolution state:

$$t \mapsto R(t) = (\ell(t), \tau(t), \Lambda(t), \epsilon(t), \Sigma(t)) \quad (7)$$

At each t , the map uses $R(t)$ inside $\Pi_{C,R(t)}$, $|r_{C,R(t)}|$, and $O(C, R(t))$. Early t : coarse (ℓ, τ, Λ) . As the process runs, $R(t)$ sharpens. t is *when* the observation is happening; $R(t)$ is *how sharply* it can distinguish at that moment.

Table 2: Components of resolution $R(t)$ at process time t .

Component	Meaning
$\ell(t)$	Spatial resolution
$\tau(t)$	Temporal sampling window
$\Lambda(t)$	Bandwidth cutoff
$\epsilon(t)$	Distinction tolerance
$\Sigma(t)$	Noise / error covariance

(ℓ, τ, Λ) improve as t advances; (ϵ, Σ) carry apparatus and environmental limits at each t . Subscripts on C, R mean evaluation at the resolution state $R(t)$ reached at process time t .

2.5 Projection operator

$$(\Pi_{C,R}\Phi)(x) = \int_U K_{C,R}(x, y) A_C[\Phi](y) d\mu_\eta(y) \quad (8)$$

where $K_{C,R}$ is a smooth sampling kernel determined by $(\ell, \tau, \Lambda, \epsilon, \Sigma)$ and weighted by context C along the throughline, and A_C is the apparatus functional.

2.6 Residual

$$r_{C,R}[\Phi] = \Phi - \iota_C \Pi_{C,R}\Phi \quad (9)$$

where ι_C embeds the projected observable back into the modelling space. The residual is **measurable**:

$$|r_{C,R}|^2 = \int_U \langle r_{C,R}(x), r_{C,R}(x) \rangle d\mu_\eta(x) \quad (10)$$

This is exactly the information a pure Minkowski projection suppresses.

3 Effective metric and interval

3.1 Metric family

$$g_{\mu\nu}^{C,R} = \eta_{\mu\nu} + h_{\mu\nu}^{C,R}[\Pi_{C,R}\Phi, r_{C,R}] \quad (11)$$

First-order calculus-compatible correction:

$$h_{\mu\nu}^{C,R} = a \partial_\mu \rho_{C,R} \partial_\nu \rho_{C,R} + b J_\mu^r J_\nu^r + c \Xi_{\mu\nu}^{C,R} \quad (12)$$

where

$$\rho_{C,R} = |\Pi_{C,R}\Phi|^2, \quad J_\mu^r = \partial_\mu |r_{C,R}|^2 \quad (13)$$

and $\Xi_{\mu\nu}^{C,R}$ is a fragmentation / decoherence tensor from projected gradients and noise.

3.2 Interval

$$ds_{C,R}^2 = g_{\mu\nu}^{C,R} dx^\mu dx^\nu = \eta_{\mu\nu} dx^\mu dx^\nu + h_{\mu\nu}^{C,R} dx^\mu dx^\nu \quad (14)$$

3.3 Minkowski recovery condition

Minkowski is recovered when

$$\Pi_{C,R}\Phi \text{ is uniform,} \quad r_{C,R} \rightarrow 0, \quad \Xi_{\mu\nu}^{C,R} \rightarrow 0. \quad (15)$$

Equivalently: **flat η is the (C, R) -low-residual limit of the projected functional—a simplification that converges, not a discovered spacetime.**

4 The refinement algorithm

CMF is not parameter tuning. It is a **process**:

1. **Declare C** —fix the environment and observer throughline (e.g. satellite position, target, ρ_{env}).
2. **Set process time t** —resolve $R(t) = (\ell(t), \tau(t), \Lambda(t), \epsilon(t), \Sigma(t))$ and evaluate at this moment of observation.
3. **Project**—compute $\Pi_{C,R(t)}\Phi$.
4. **Measure residual**—compute $|r_{C,R(t)}|^2$.
5. **Advance process time**—as t increases, $R(t)$ sharpens and the functional refines.
6. **Recover Minkowski**—when residual falls below tolerance, flat geometry is a valid zeroth-order description.

$$\boxed{(\mathcal{M}, \eta) \longrightarrow (\mathcal{M}, \eta, \Phi, \Pi_{C,R}, r_{C,R}, g^{C,R})} \quad (16)$$

5 Observers and events

Events are **not** primitive points. An observation is

$$O_i(C, R) = \int_U K_{C,R}^i(x) F_i(\Phi, \partial\Phi, g^{C,R}) d\mu_{g^{C,R}} + \epsilon_i \quad (17)$$

a finite-resolution projection of field-process into an observable, with sampling noise ϵ_i .

Table 3: Standard practice vs. CMF.

Standard	CMF
Fixed $F[\Phi]$	Process functional $F[\Phi; C, R(t)]$
C, R bolted on as inputs	C, R inside the map
Same answer at all times	Functional sharpens with process time
Event = point on $\mathcal{M}$	Event = kernel projection

6 Dynamics (optional extension)

Field dynamics remain calculus-based. Action:

$$S[\Phi] = \int_{\mathcal{M}} \left[\frac{1}{2} \eta^{\mu\nu} \langle \partial_\mu \Phi, \partial_\nu \Phi \rangle - V(\rho_{C,R}, \Xi) + \frac{\alpha}{2} |r_{C,R}|^2 \right] d^4x \quad (18)$$

Euler–Lagrange equations yield PDEs with residual source terms:

$$\square_\eta \Phi + \frac{\partial V}{\partial \Phi} = J_r \quad (19)$$

The residual source J_r is what old Minkowski projection would drop.

7 When the utilitarian simplification suffices

Use flat η when the accounting simplification is adequate for the task—not because spacetime is true.

Table 4: When flat η suffices vs. when to carry CMF.

Situation	Recommendation
Back-of-envelope kinematics	Flat η —simplification suffices
Lab frame, long integration, buried error budget	Often suffices
Satellite line-of-sight through density gradient	Carry CMF
Early integration, targeting decisions	Carry CMF
Residual-dominated regime	Do not reify η as ontology

Flat Minkowski is a **utilitarian simplification**, not relational reality. CMF adds no overhead when the simplification is adequate: set $h \rightarrow 0$ and recover η . Carry C and $R(t)$ when the accounting task requires what the projection suppresses.

8 Stipulations vs. determination

CMF is a **cogenetic** projection formalism [?]: observables co-emerge with context and process time. Judging it against general relativity or quantum field theory as a “correction” is the wrong frame—those theories already presuppose flat η as stage [?]. The correct comparison is **CMF vs. standard Minkowski practice**.

The distinction is methodological, not a claim that CMF never predicts anything:

- **Minkowski stipulates**—at a given setup, standard practice *fixes* $|r_{C,R}| = 0$, $g = \eta$, and a functional with no R -map inside it. These are prior commitments of the calculus stage, not outputs of an observation process.

- **CMF determines**—at each process time t , the map resolves $R(t)$ and computes $|r_{C,R(t)}|$, $g^{C,R(t)} = \eta + h$, and $O(C, R(t))$ from (Φ, C, t) via project $\rightarrow$ measure residual $\rightarrow$ evaluate observer (??).
- **CMF predicts**—over process time, the formalism claims refinement behaviour: falling residual, sharpening $O(C, R)$, and a bolted-on $F[\Phi]$ that cannot track any of it.

8.1 What Minkowski stipulates

Table 5: Values fixed by standard Minkowski practice before the observation process runs.

Stipulation	Value
Residual	$ r_{C,R} = 0$
Metric	$g^{\mu\nu} = \eta^{\mu\nu}$
Interval correction	$ g_{xx} - \eta_{xx} = 0$
Functional	$F[\Phi]$ with no R -dependence
Over process time	No change— C and R bolted on externally

These are utilitarian defaults—often adequate. They are not computed from the throughline environment or the running observation.

8.2 What CMF determines at each R

At every process time, CMF **derives** quantities Minkowski practice cannot represent:

- **Residual** $|r_{C,R}|^2 = |\Phi - \iota_C \Pi_{C,R} \Phi|^2$ along the throughline.
- **Metric correction** $h[\Pi\Phi, r]$ and $g^{C,R} = \eta + h$; in vacuum $g_{xx} = 1$, where $|C| > 0$ and the process field has structure, $|g_{xx} - \eta_{xx}|$ is generally nonzero.
- **Observer reading** $O(C, R)$ from (??)—a kernel projection of $\rho_{C,R} = |\Pi_{C,R} \Phi|^2$ under $d\mu_{g^{C,R}}$, not raw $|\Phi|^2$ at a point.

What we do *not* compare. CMF does **not** claim that $O(C, R)$ converges to raw $|\Phi|^2$ at a point. Comparing $O(C, R)$ to $|\Phi(0, 0)|^2$ is the wrong metric. The correct comparison is **stipulation vs. determination**: what Minkowski *fixes in advance* versus what CMF *computes from the map* at each R .

8.3 What CMF predicts over R

Table 6: Process behaviour: stipulation vs. determination vs. prediction.

Minkowski (stipulates)	CMF (determines / predicts)
$ r_{C,R} = 0$ always	determines $ r_{C,R} ^2$; coarse-to-fine refinement drives it down
O independent of R	determines $O(C, R(t))$; predicts sharpening with R
$g^{\mu\nu} = \eta^{\mu\nu}$	determines $h[\Pi\Phi, r] \neq 0$ where $ C > 0$
$ g_{xx} - \eta_{xx} = 0$	determines $ g_{xx} - \eta_{xx} $ at target
No R -map inside $F[\Phi]$	predicts $ O(C, R) - O(C, R_{\text{late}}) \rightarrow 0$ as process runs

Testable claims.

- P1: Residual refinement** (*predicts*). For fixed C and Φ , coarse-to-fine process time drives $|r_{C,R}|^2$ down over the main observation trajectory. Very late R may show a local uptick when projection sharpens on a structured target—the claim is refinement, not strict monotonicity at every step.
- P2: Functional dependence on R** (*predicts*). $F[\Phi; C, R(t)]$ and $O(C, R)$ vary with R ; a bolted-on $F[\Phi] = \int |\Phi|^2 d\mu$ does not.
- P3: Context-driven metric correction** (*determines*). $h[\Pi\Phi, r]$ correlates with throughline context $|C|$, not identically zero as η stipulates.
- P4: Interval correction** (*determines*). $|g_{xx} - \eta_{xx}|$ at the target is computed by CMF; Minkowski stipulates zero. In vacuum, $g_{xx} = 1$ throughout; at the structured target the correction need not shrink with R —the point is that CMF *represents* it.
- P5: Process refinement** (*predicts*). $|O(C, R) - O(C, R_{\text{late}})|$ decreases with R . Standard Minkowski practice offers no R -dependence inside $F[\Phi]$; a bolted-on fixed integral is unchanged at all process times.

Verification (satellite demo). The reference implementation separates the three roles on the worked setup below. **Stipulations:** Minkowski fixes $|r| = 0$, $|g_{xx} - 1| = 0$, and $F[\Phi] \approx 4.52$ at all R . **Determinations:** at early R , CMF computes $|r_{C,R}|^2 \approx 3.1$ and $|g_{xx} - 1| \approx 0.06$; at late R , $|r_{C,R}|^2 \approx 0.05$ and $|g_{xx} - 1| \approx 2.4$ at the target—nonzero structure Minkowski cannot represent. **Predictions:** residual falls from early to late R (with a modest uptick in the final refinement steps on this grid); $|O(C, R) - O(C, R_{\text{late}})|$ approaches zero; the bolted-on functional is unchanged throughout.

9 Worked example: satellite point detector

Setup.

- Observer (satellite) at altitude $x = -7$, target at surface $x = 0$.
- Environmental density: $\rho_{\text{env}}(x) = \exp(-|x|/H)$ (exponential atmosphere).
- Context: $C = \hat{n} \cdot \nabla \rho_{\text{env}}$ along the throughline (??).
- Process field Φ : Gaussian target packet on the surface.

Results (reference implementation). At early process time ($R = t = -4$), projection is coarse: $|r_{C,R}|^2 \approx 3.1$ and $g_{xx} \approx 1.06$ at the target (cf. Minkowski $\eta_{xx} = 1$). At late process time ($R = t = 4$), the residual falls to $|r_{C,R}|^2 \approx 0.05$; refinement error $|O(C, R) - O(C, R_{\text{late}})|$ approaches zero. The process functional $F[\Phi; C, R(t)]$ shifts from ≈ 0.002 to ≈ 0.58 while a fixed functional $F[\Phi] = \int |\Phi|^2 d\mu \approx 4.52$ is unchanged at all R — C and R play no role inside the bolted-on map. Vacuum regions carry $g_{xx} = 1$ throughout; flat Minkowski is the valid zeroth-order description where $|r_{C,R}|$ is negligible, not necessarily where the resolved process field has structure.

10 Reference implementation

11 Relation to deeper process geometry

The *Cogenesis* construction [?] gives a deeper form:

$$d\Sigma_{C,R}^2 = \langle \delta_{C,R}\Phi, G_{C,R}[\Phi] \delta_{C,R}\Phi \rangle + r_\Sigma \quad (20)$$




Figure 1: Stipulations vs. determination vs. prediction (python demo.py). Top row: CMF *determines* residual and interval correction; Minkowski *stipulates* both zero. Bottom row: CMF *predicts* refinement and R -dependence; Minkowski *stipulates* no R -map.




Figure 2: Reference implementation figures (python demo.py).

Table 7: Release artefacts.

File	Role
<code>cmf.tex</code>	LaTeX source (this document)
<code>cmf.pdf</code>	Compiled release PDF
<code>model.py</code>	Core CMF: C as throughline gradient; $t \mapsto R(t)$ resolution state
<code>demo.py</code>	Satellite demonstration; generates release figures
<code>requirements.txt</code>	<code>numpy</code> , <code>scipy</code> , <code>matplotlib</code>
<code>README.md</code> , <code>LICENSE</code>	Release metadata
<code>output/</code>	Demonstration figures

with classical metric geometry recovered by projection:

$$ds^2 = \pi_{\text{geom}}(d\Sigma_{C,R}^2) \quad (21)$$

CMF represents this projection as an effective metric family $g_{\mu\nu}^{C,R}$ —calculus-safe, without pretending the full $d\Sigma^2$ has been captured by ds^2 alone. Where a single shared causal cone and metric already unify sectors at the calculus level [?], CMF specifies how flat η arises as the low-residual projection of that deeper process under declared context and process time.

12 Conclusion

CMF keeps the Minkowski projection calculus without reifying it—against the process calculus that carries C , $R(t)$, and residual:

- **Same machinery**—manifolds, tensors, kernels, PDEs, variational methods.
- **Same simplification exit**—flat η when $|r_{C,R}|$ is negligible and the accounting task allows it.
- **No reification**— C and $R(t)$ inside the functional; process time indexes resolution, not a fused spacetime ontology.
- **Relational dynamics**—events, intervals, and metrics from observation process, not primitive points on $\mathcal{M}$.

Reified:	$(\mathcal{M}, \eta_{\mu\nu})$	as primitive spacetime
CMF:	$(\mathcal{M}, \eta_{\mu\nu}, \Phi, \Pi_{C,R}, r_{C,R}, g_{\mu\nu}^{C,R})$	flat η when $ r_{C,R} \rightarrow 0$ —utilitarian simplification, not ontology.

References

- [1] M. Petrik and Militant.AI, *Cogenesis*, Zenodo, 2026. [doi:10.5281/zenodo.20539999](https://doi.org/10.5281/zenodo.20539999)
- [2] M. Petrik and Militant.AI, *One-Metric Quantum Gravity: Unifying SR, GR, and QM with a Shared Causal Cone*, Zenodo, 2025. [doi:10.5281/zenodo.17731104](https://doi.org/10.5281/zenodo.17731104)